

QR Exercise 3.5 MATH531

Jared Andreatta

2025-01-28

Getting started

First here is an example dataset that you can use to illustrate OLS computations. Loading the AudiA4 data. The goal is to be able to predict the price of a used A4 based on the mileage and model year.

```
library( fields) # load fields package

## Warning: package 'fields' was built under R version 4.4.2

## Loading required package: spam

## Warning: package 'spam' was built under R version 4.4.2

## Spam version 2.11-1 (2025-01-20) is loaded.
## Type 'help( Spam)' or 'demo( spam)' for a short introduction
## and overview of this package.
## Help for individual functions is also obtained by adding the
## suffix '.spam' to the function name, e.g. 'help( chol.spam)'.

##
## Attaching package: 'spam'

## The following objects are masked from 'package:base':
##
##      backsolve, forwardsolve

## Loading required package: viridisLite

##
## Try help(fields) to get started.
```

```
load("AudiA4.rda" )
head( AudiA4)
```

```
##      year price mileage distance
## 58  2020 39649    3848         29
## 145 2020 43175    3962          7
## 10   2020 43675    5316          7
## 52   2020 40649    5417         29
## 143 2020 42175    5846          7
## 9    2020 45675    6539          7
```

```
# scale and adjust variables to be easier to interpret.
year<- AudiA4$year- min(AudiA4$year)
mileage<- AudiA4$mileage/1000
Y<- AudiA4$price/1000

X<- cbind( 1, mileage, year )

fit<- lm( Y ~ X - 1)
```

For those students rusty with linear algebra in R here is a digression to find OLS estimates found “by hand”.

```
betaHat<- solve(t(X)%*%X)%*%t(X)%*%Y
# check
bothBetas<- cbind(betaHat,fit$coefficients )
print( bothBetas)
```

```
##           [,1]      [,2]
##      14.3751275 14.3751275
## mileage -0.1241398 -0.1241398
## year      1.1550828  1.1550828
```

```
XbetaHat<- X%*%betaHat
test.for.zero(XbetaHat, fit$fitted.values )
```

```
## PASSED test at tolerance 1e-08
```

Problem 1

- Define the QR decomposition of an $n \times p$ matrix X .
Distinguish between the *skinny* QR where Q is $n \times p$ and R is $p \times p$ and the *fat* one where where Q is $n \times n$ and R is $n \times p$. (The skinny one is usually what is computed.)
 - The “skinny” QR Decomposition refers to the situation in which $Q \in \mathbb{R}^{n \times p}$ where Q is comprised of orthonormal column vectors, and $n > p$, and more often than not, $n \gg p$, as there tends to be more observations than features in the dataset; $R \in \mathbb{R}^{n \times n}$ is an upper right triangular matrix. In this case, the columns of X are expressed as a linear combination with the orthonormal columns of Q .
The “fat” QR Decomposition, where $Q \in \mathbb{R}^{n \times n}$ and $R \in \mathbb{R}^{n \times p}$. The difference in this is the number of columns in Q ; Q has n columns that form an orthonormal basis of \mathbb{R}^n . R is a rectangular matrix of dimension $n \times p$, where the upper triangular block has dimension of $p \times p$.
To summarize, the “skinny” decomposition is often used for “tall” matrices due to its computational properties, whereas the “fat” decomposition is used to form an orthonormal basis that spans \mathbb{R}^n .
- Who invented the QR decomposition?
 - The QR algorithm was first proposed by John G. F. Francis and Vera N. Kublanovskaya (not collaboratively) in the late 1950s.
- Using the function `qr` and related functions (`qr.coef`, `qr.Q`, etc.) in R find the skinny QR decomposition of X (List the first 5 rows of “Q” and all of “R”)

```
q = qr(X)
head(qr.Q(q))
```

```
##           [,1]      [,2]      [,3]
## [1,] -0.09166985 -0.1191055 -0.005470635
## [2,] -0.09166985 -0.1188379 -0.005921018
## [3,] -0.09166985 -0.1156595 -0.011270299
## [4,] -0.09166985 -0.1154224 -0.011669323
## [5,] -0.09166985 -0.1144153 -0.013364184
## [6,] -0.09166985 -0.1127886 -0.016102036
```

```
print(qr.R(q))
```

```
##           mileage      year
## [1,] -10.90871 -595.4751 -137.7798
## [2,]  0.00000  426.0015  -35.7134
## [3,]  0.00000   0.0000  -21.2199
```

Problem 2

Use the QR decomposition to compute the OLS estimate for the Audi A4 data and verify that it is same as the `lm` result above.

```
QR_X <- qr.Q(q) %*% qr.R(q)
QRbetaHat <- QR_X %*% betaHat
test.for.zero(XbetaHat, fit$fitted.values )
```

```
## PASSED test at tolerance 1e-08
```

```
head(QRbetaHat)
```

```
##           [,1]
## [1,] 33.53385
## [2,] 33.51969
## [3,] 33.35161
## [4,] 33.33907
## [5,] 33.28581
## [6,] 33.19979
```

```
head(XbetaHat)
```

```
##           [,1]
## [1,] 33.53385
## [2,] 33.51969
## [3,] 33.35161
## [4,] 33.33907
## [5,] 33.28581
## [6,] 33.19979
```

Both estimators give the same output