

R Lab 4

Jared Andreatta

2025-04-07

6.5.1 Subset Selection Methods

Best Subset Selection

```
library(ISLR2)
```

```
## Warning: package 'ISLR2' was built under R version 4.4.3
```

```
names(Hitters)
```

```
## [1] "AtBat"      "Hits"       "HmRun"      "Runs"       "RBI"        "Walks"
## [7] "Years"      "CAtBat"     "CHits"      "CHmRun"     "CRuns"      "CRBI"
## [13] "CWalks"     "League"     "Division"   "PutOuts"    "Assists"    "Errors"
## [19] "Salary"     "NewLeague"
```

```
dim(Hitters) # Dimension of data matrix
```

```
## [1] 322 20
```

```
sum(is.na(Hitters$Salary)) # Sums up NA values for salary of hitters
```

```
## [1] 59
```

```
Hitters <- na.omit(Hitters) # Omitting observations where any variable is NA
sum(is.na(Hitters))
```

```
## [1] 0
```

We can use the **leaps** library.

```
# install.packages("leaps")
library(leaps)
```

```
## Warning: package 'leaps' was built under R version 4.4.3
```

```
# regsubsets() automatically does best subset selection using RSS
regfit.full <- regsubsets(Salary ~ ., data=Hitters)
summary(regfit.full)
```

```
## Subset selection object
## Call: regsubsets.formula(Salary ~ ., data = Hitters)
## 19 Variables (and intercept)
##           Forced in Forced out
## AtBat      FALSE      FALSE
## Hits       FALSE      FALSE
## HmRun       FALSE      FALSE
## Runs       FALSE      FALSE
## RBI        FALSE      FALSE
## Walks      FALSE      FALSE
## Years      FALSE      FALSE
## CAtBat     FALSE      FALSE
## CHits      FALSE      FALSE
## CHmRun     FALSE      FALSE
## CRuns      FALSE      FALSE
## CRBI       FALSE      FALSE
## CWalks     FALSE      FALSE
## LeagueN    FALSE      FALSE
## DivisionW  FALSE      FALSE
## PutOuts    FALSE      FALSE
## Assists    FALSE      FALSE
## Errors     FALSE      FALSE
## NewLeagueN FALSE      FALSE
## 1 subsets of each size up to 8
## Selection Algorithm: exhaustive
##           AtBat Hits HmRun Runs RBI Walks Years CAtBat CHits CHmRun CRuns CRBI
## 1 ( 1 ) " " " " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 2 ( 1 ) " " "*" " " " " " " " " " " " " " " " " " " " " " "
## 3 ( 1 ) " " "*" " " " " " " " " " " " " " " " " " " " " "
## 4 ( 1 ) " " "*" " " " " " " " " " " " " " " " " " " " " "
## 5 ( 1 ) "*" "*" " " " " " " " " " " " " " " " " " " " " "
## 6 ( 1 ) "*" "*" " " " " " " "*" " " " " " " " " " " " " " "
## 7 ( 1 ) " " "*" " " " " " " "*" " " "*" "*" "*" " " " " "
## 8 ( 1 ) "*" "*" " " " " " " "*" " " " " "*" "*" "*" " " "
##           CWalks LeagueN DivisionW PutOuts Assists Errors NewLeagueN
## 1 ( 1 ) " " " " " " " " " " " " " "
## 2 ( 1 ) " " " " " " " " " " " "
## 3 ( 1 ) " " " " " " "*" " " " " "
## 4 ( 1 ) " " " " "*" "*" " " " " "
## 5 ( 1 ) " " " " "*" "*" " " " " "
## 6 ( 1 ) " " " " "*" "*" " " " " "
## 7 ( 1 ) " " " " "*" "*" " " " " "
## 8 ( 1 ) "*" " " "*" "*" " " " " " "
```

Asterisks indicate that the variable is included in the corresponding model. The **nvmax** parameter in the function lets us predetermine the max amount of variables.

```
regfit.full <- regsubsets(Salary ~ ., data = Hitters, nvmax = 19) # 19 vars
reg.summary <- summary(regfit.full)
```



```
# Columns of summary object
names(reg.summary)
```

```
## [1] "which" "rsq" "rss" "adjr2" "cp" "bic" "outmat" "obj"
```

```
# R^2
reg.summary$rsq
```

```
## [1] 0.3214501 0.4252237 0.4514294 0.4754067 0.4908036 0.5087146 0.5141227
## [8] 0.5285569 0.5346124 0.5404950 0.5426153 0.5436302 0.5444570 0.5452164
## [15] 0.5454692 0.5457656 0.5459518 0.5460945 0.5461159
```

```
# Plotting the subset RSS and adjusted R^2
par(mfrow = c(1, 2))
plot(reg.summary$rss, xlab = "Number of Variables",
     ylab = "RSS", type = "l")
plot(reg.summary$adjr2, xlab = "Number of Variables",
     ylab = "Adjusted RSq", type = "l")
```



```
# Finding max adj R2
which.max(reg.summary$adjr2) # Max # of vars
```



```
## [1] 11
```

```
plot(reg.summary$adjr2, xlab = "Number of Variables", ylab = "Adjusted RSq", type = "l") # Plotting adj
points(11, reg.summary$adjr2[11], col = "orange", cex = 2, pch = 20) # Plotting point on line
```

```
# Analogous procedure for Cp
plot(reg.summary$cp, xlab = "Number of Variables", ylab = "Cp", type = "l")
which.min(reg.summary$cp)
```

```
## [1] 10
```


```
points(10, reg.summary$cp[10], col = "red", cex = 2, pch = 20)
```



```
# Analogous procedure for BIC
plot(reg.summary$bic, xlab = "Number of Variables", ylab = "BIC", type = "l")
which.min(reg.summary$bic)
```


```
## [1] 6
```

```
points(6, reg.summary$bic[6], col = "red", cex = 2, pch = 20)
```




The `regsubsets()` function also has a built in plot command which can display selected variables for a model according to different metrics. The top row contains a black square for a selected variable in the optimal model


```
plot(regfit.full, scale = "r2") # R2
```




```
plot(regfit.full, scale = "adjr2") # Adj R^2
```



```
plot(regfit.full, scale = "Cp") # C_p
```



```
plot(regfit.full, scale = "bic") # BIC
```

```
# Coefficients of 6th model
coef(regfit.full, 6)
```

```
## (Intercept)      AtBat      Hits      Walks      CRBI      DivisionW
##  91.5117981    -1.8685892    7.6043976    3.6976468    0.6430169   -122.9515338
##      PutOuts
##    0.2643076
```

Forward and Backward Stepwise Selection

We can also use `regsubsets()` to perform stepwise selection using the ‘method’ parameter. Note that compared to the full best subset selection, the best models for 1-6 variables are all the same, so we will look at a 7 variable model.

```
# Forward
regfit.fwd <- regsubsets(Salary ~ ., data = Hitters, nvmax = 19, method = "forward")
summary(regfit.fwd)
```

```
## Subset selection object
## Call: regsubsets.formula(Salary ~ ., data = Hitters, nvmax = 19, method = "forward")
## 19 Variables (and intercept)
##      Forced in Forced out
## AtBat      FALSE      FALSE
## Hits       FALSE      FALSE
```

```

## HmRun          FALSE      FALSE
## Runs           FALSE      FALSE
## RBI            FALSE      FALSE
## Walks          FALSE      FALSE
## Years          FALSE      FALSE
## CAtBat         FALSE      FALSE
## CHits          FALSE      FALSE
## CHmRun         FALSE      FALSE
## CRuns          FALSE      FALSE
## CRBI           FALSE      FALSE
## CWalks         FALSE      FALSE
## LeagueN        FALSE      FALSE
## DivisionW      FALSE      FALSE
## PutOuts        FALSE      FALSE
## Assists        FALSE      FALSE
## Errors         FALSE      FALSE
## NewLeagueN     FALSE      FALSE
## 1 subsets of each size up to 19
## Selection Algorithm: forward
##           AtBat Hits HmRun Runs RBI Walks Years CAtBat CHits CHmRun CRuns CRBI
## 1 ( 1 ) " " " " " " " " " " " " " " " " " " " "
## 2 ( 1 ) " " "*" " " " " " " " " " " " " " " " "
## 3 ( 1 ) " " "*" " " " " " " " " " " " " " " " "
## 4 ( 1 ) " " "*" " " " " " " " " " " " " " " " "
## 5 ( 1 ) "*" "*" " " " " " " " " " " " " " " " "
## 6 ( 1 ) "*" "*" " " " " " " "*" " " " " " " " " " "
## 7 ( 1 ) "*" "*" " " " " " " "*" " " " " " " " " " "
## 8 ( 1 ) "*" "*" " " " " " " "*" " " " " " " "*" " "
## 9 ( 1 ) "*" "*" " " " " " " "*" " " "*" " " " " "*" "
## 10 ( 1 ) "*" "*" " " " " " " "*" " " "*" " " " " "*" "
## 11 ( 1 ) "*" "*" " " " " " " "*" " " "*" " " " " "*" "
## 12 ( 1 ) "*" "*" " " " " "*" " " " "*" " " " " "*" "
## 13 ( 1 ) "*" "*" " " " " "*" " " " "*" " " " " "*" "
## 14 ( 1 ) "*" "*" "*" "*" " " " " "*" " " " " " "*" "
## 15 ( 1 ) "*" "*" "*" "*" " " " " "*" " "*" " " " "*" "
## 16 ( 1 ) "*" "*" "*" "*" "*" "*" " " " " "*" " " " "*" "
## 17 ( 1 ) "*" "*" "*" "*" "*" "*" " " " " "*" " " " "*" "
## 18 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "*" " " "*" " " " "*" "
## 19 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" " " "*" "
##           CWalks LeagueN DivisionW PutOuts Assists Errors NewLeagueN
## 1 ( 1 ) " " " " " " " " " " " "
## 2 ( 1 ) " " " " " " " " " " " "
## 3 ( 1 ) " " " " " " "*" " " " " "
## 4 ( 1 ) " " " " "*" "*" " " " "
## 5 ( 1 ) " " " " "*" "*" " " " "
## 6 ( 1 ) " " " " "*" "*" " " " "
## 7 ( 1 ) "*" " " " "*" "*" " " " "
## 8 ( 1 ) "*" " " " "*" "*" " " " "
## 9 ( 1 ) "*" " " " "*" "*" " " " "
## 10 ( 1 ) "*" " " " "*" "*" "*" " " "
## 11 ( 1 ) "*" "*" " " "*" "*" " " " "
## 12 ( 1 ) "*" "*" " " "*" "*" " " " "
## 13 ( 1 ) "*" "*" " " "*" "*" "*" " "
## 14 ( 1 ) "*" "*" " " "*" "*" "*" " "

```

```
## 15 ( 1 ) "*" "*" "*" "*" "*" "*" " "
## 16 ( 1 ) "*" "*" "*" "*" "*" "*" " "
## 17 ( 1 ) "*" "*" "*" "*" "*" "*" "*"
## 18 ( 1 ) "*" "*" "*" "*" "*" "*" "*"
## 19 ( 1 ) "*" "*" "*" "*" "*" "*" "*"

```

Backward

```
regfit.bwd <-regsubsets(Salary ~ ., data = Hitters, nvmax = 19, method = "backward")
summary(regfit.bwd)

```

```
## Subset selection object
## Call: regsubsets.formula(Salary ~ ., data = Hitters, nvmax = 19, method = "backward")
## 19 Variables (and intercept)
##              Forced in Forced out
## AtBat          FALSE          FALSE
## Hits           FALSE          FALSE
## HmRun          FALSE          FALSE
## Runs           FALSE          FALSE
## RBI            FALSE          FALSE
## Walks          FALSE          FALSE
## Years          FALSE          FALSE
## CAtBat         FALSE          FALSE
## CHits          FALSE          FALSE
## CHmRun         FALSE          FALSE
## CRuns          FALSE          FALSE
## CRBI           FALSE          FALSE
## CWalks         FALSE          FALSE
## LeagueN        FALSE          FALSE
## DivisionW      FALSE          FALSE
## PutOuts        FALSE          FALSE
## Assists        FALSE          FALSE
## Errors         FALSE          FALSE
## NewLeagueN     FALSE          FALSE
## 1 subsets of each size up to 19
## Selection Algorithm: backward
##              AtBat Hits HmRun Runs RBI Walks Years CAtBat CHits CHmRun CRuns CRBI
## 1 ( 1 ) " " " " " " " " " " " " " " " " "*" " "
## 2 ( 1 ) " " "*" " " " " " " " " " " " " "*" " "
## 3 ( 1 ) " " "*" " " " " " " " " " " " " "*" " "
## 4 ( 1 ) "*" "*" " " " " " " " " " " " " "*" " "
## 5 ( 1 ) "*" "*" " " " " " " "*" " " " " " " "*" " "
## 6 ( 1 ) "*" "*" " " " " " " "*" " " " " " " "*" " "
## 7 ( 1 ) "*" "*" " " " " " " "*" " " " " " " "*" " "
## 8 ( 1 ) "*" "*" " " " " " " "*" " " " " " " "*" "*"
## 9 ( 1 ) "*" "*" " " " " " " "*" " " "*" " " " " "*" "*"
## 10 ( 1 ) "*" "*" " " " " " " "*" " " "*" " " " " "*" "*"
## 11 ( 1 ) "*" "*" " " " " " " "*" " " "*" " " " " "*" "*"
## 12 ( 1 ) "*" "*" " " "*" " " "*" " " "*" " " " " "*" "*"
## 13 ( 1 ) "*" "*" " " "*" " " "*" " " "*" " " " " "*" "*"
## 14 ( 1 ) "*" "*" "*" "*" " " "*" " " "*" " " " " "*" "*"
## 15 ( 1 ) "*" "*" "*" "*" " " "*" " " "*" "*" " " " " "*" "*"
## 16 ( 1 ) "*" "*" "*" "*" "*" "*" " " "*" "*" " " " " "*" "*"
## 17 ( 1 ) "*" "*" "*" "*" "*" "*" " " "*" "*" " " " " "*" "*"
## 18 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "*" "*" " " " " "*" "*"

```

```
## 19 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*"
##           CWalks LeagueN DivisionW PutOuts Assists Errors NewLeagueN
## 1 ( 1 ) " " " " " " " " " " " " " "
## 2 ( 1 ) " " " " " " " " " " " " " "
## 3 ( 1 ) " " " " " " "*" " " " " " "
## 4 ( 1 ) " " " " " " "*" " " " " " "
## 5 ( 1 ) " " " " " " "*" " " " " " "
## 6 ( 1 ) " " " " "*" "*" " " " " " "
## 7 ( 1 ) "*" " " "*" "*" " " " " " "
## 8 ( 1 ) "*" " " "*" "*" " " " " " "
## 9 ( 1 ) "*" " " "*" "*" " " " " " "
## 10 ( 1 ) "*" " " "*" "*" "*" " " " " "
## 11 ( 1 ) "*" "*" "*" "*" "*" " " " " "
## 12 ( 1 ) "*" "*" "*" "*" "*" " " " " "
## 13 ( 1 ) "*" "*" "*" "*" "*" "*" " " " "
## 14 ( 1 ) "*" "*" "*" "*" "*" "*" " " " "
## 15 ( 1 ) "*" "*" "*" "*" "*" "*" " " " "
## 16 ( 1 ) "*" "*" "*" "*" "*" "*" " " " "
## 17 ( 1 ) "*" "*" "*" "*" "*" "*" "*" " " "
## 18 ( 1 ) "*" "*" "*" "*" "*" "*" "*" " " "
## 19 ( 1 ) "*" "*" "*" "*" "*" "*" "*" " " "
```

```
coef(regfit.full, 7)
```

```
## (Intercept)           Hits           Walks           CAtBat           CHits           CHmRun
## 79.4509472    1.2833513    3.2274264    -0.3752350    1.4957073    1.4420538
## DivisionW      PutOuts
## -129.9866432    0.2366813
```

```
coef(regfit.fwd, 7)
```

```
## (Intercept)           AtBat           Hits           Walks           CRBI           CWalks
## 109.7873062   -1.9588851    7.4498772    4.9131401    0.8537622   -0.3053070
## DivisionW      PutOuts
## -127.1223928    0.2533404
```

```
coef(regfit.bwd, 7)
```

```
## (Intercept)           AtBat           Hits           Walks           CRuns           CWalks
## 105.6487488   -1.9762838    6.7574914    6.0558691    1.1293095   -0.7163346
## DivisionW      PutOuts
## -116.1692169    0.3028847
```

Choosing Among Models Using the Validation-Set Approach and Cross-Validation

We start with the validation set approach.

```

set.seed(1)
# Create random vector of elements of TRUE or FALSE if the corresponding observation is in the training
train <- sample(c(TRUE, FALSE), nrow(Hitters), replace = TRUE)
test <- (!train)

# Apply regsubsets() to training data
regfit.best <- regsubsets(Salary ~ ., data = Hitters[train, ], nvmax = 19)

# Make test data matrix for model
test.mat <- model.matrix(Salary ~ ., data = Hitters[test, ])

val.errors <- rep(NA, 19)
# Using test matrix to compute test MSE
for (i in 1:19) {
  coefi <- coef(regfit.best, id = i)
  pred <- test.mat[, names(coefi)] %*% coefi
  val.errors[i] <- mean((Hitters$Salary[test] - pred)^2)
}

val.errors

## [1] 164377.3 144405.5 152175.7 145198.4 137902.1 139175.7 126849.0 136191.4
## [9] 132889.6 135434.9 136963.3 140694.9 140690.9 141951.2 141508.2 142164.4
## [17] 141767.4 142339.6 142238.2

```

```
which.min(val.errors) # Min errors for 7 vars
```

```
## [1] 7
```

```
coef(regfit.best, 7) # Coefs for 7 var model
```

```
## (Intercept)      AtBat      Hits      Walks      CRuns      CWalks
## 67.1085369 -2.1462987  7.0149547  8.0716640  1.2425113 -0.8337844
## DivisionW      PutOuts
## -118.4364998  0.2526925
```

```
# This function just mimics the code above
```

```

predict.regsubsets <- function(object, newdata, id, ...) {
  form <- as.formula(object$call[[2]])
  mat <- model.matrix(form, newdata)
  coefi <- coef(object, id = id)
  xvars <- names(coefi)
  mat[, xvars] %*% coefi
}

```

```
## Now we apply the regsubsets() to the full data set ##
```

```
regfit.best <- regsubsets(Salary ~ ., data = Hitters, nvmax = 19)
```

```
# NOTE: the subset of variables is different than the set of variables from doing this on the training
```

```
coef(regfit.best, 7)
```

```
## (Intercept)      Hits      Walks      CAtBat      CHits      CHmRun
```

```
##      79.4509472      1.2833513      3.2274264      -0.3752350      1.4957073      1.4420538
##      DivisionW      PutOuts
## -129.9866432      0.2366813
```

Now, we move on to the cross-validation approach.


```
# Set up for k-fold cross validation
k <- 10
n <- nrow(Hitters)
set.seed(1)
folds <- sample(rep(1:k, length = n))
cv.errors <- matrix(NA, k, 19, dimnames = list(NULL, paste(1:19)))

# For loop to perform cross validation
for (j in 1:k) {
  best.fit <- regsubsets(Salary ~. , data = Hitters[folds != j, ], nvmax = 19)
  for (i in 1:19) {
    pred <- predict(best.fit, Hitters[folds == j, ], id = i)
    cv.errors[j, i] <- mean((Hitters$Salary[folds == j] - pred)^2)
  }
}

# Using apply() to average over the columns of the matrix
mean.cv.errors <- apply(cv.errors, 2, mean)
mean.cv.errors # This approach uses the 10 variable model
```

```
##          1          2          3          4          5          6          7          8
## 143439.8 126817.0 134214.2 131782.9 130765.6 120382.9 121443.1 114363.7
##          9         10         11         12         13         14         15         16
## 115163.1 109366.0 112738.5 113616.5 115557.6 115853.3 115630.6 116050.0
##         17         18         19
## 116117.0 116419.3 116299.1
```

```
# Plotting errors. Clear minimizer at 10
par(mfrow=c(1,1))
plot(mean.cv.errors, type = "b")
```



```
# Now doing best subset selection for 10 vars
reg.best <- regsubsets(Salary~., data = Hitters, nvmax = 19)
coef(reg.best, 10)
```

```
## (Intercept)      AtBat      Hits      Walks      CAtBat      CRuns
## 162.5354420   -2.1686501    6.9180175    5.7732246   -0.1300798    1.4082490
##      CRBI      CWalks  DivisionW      PutOuts      Assists
##    0.7743122   -0.8308264  -112.3800575    0.2973726    0.2831680
```